



Docker: Building, deploying, and managing containers

Training Material

Table of Contents

Chapter 1: Introduction to Docker.....	3
1.1 What is Docker?.....	3
1.2 Understanding Containerization.....	3
1.3 Key Docker Concepts.....	4
Chapter 2: Installing and Setting Up Docker.....	5
2.1 Installing Docker.....	5
2.2 Basic Docker Commands.....	6
2.3 Configuring Docker Environment.....	6
Chapter 3: Building Docker Images.....	7
3.1 Creating Dockerfiles.....	8
3.2 Building Images.....	7
3.3 Best Practices for Image Optimization.....	8
Chapter 4: Managing Docker Containers.....	8
4.1 Running Containers.....	8
4.2 Networking in Docker.....	9
4.3 Volume Management.....	10
Chapter 5: Docker Compose for Multi-Container Applications.....	10
5.1 Introduction to Docker Compose.....	10
5.2 Defining Multi-Container Applications.....	11
5.3 Managing and Scaling Applications.....	11
Chapter 6: Deploying Docker Containers.....	12
6.1 Deployment Strategies.....	12
6.2 Using Docker Swarm.....	12
6.3 Using Kubernetes with Docker.....	13

Chapter 7: Monitoring and Logging.....	14
7.1 Monitoring Docker Containers.....	14
7.2 Logging in Docker.....	14
7.3 Centralized Logging Solutions.....	15
Chapter 8: Security Best Practices for Docker.....	16
8.1 Understanding Docker Security Risks.....	16
8.2 Implementing Security Measures.....	16
8.3 Compliance and Governance.....	16
Summary and Key Takeaways.....	16

Chapter 1: Introduction to Docker

1.1 What is Docker?

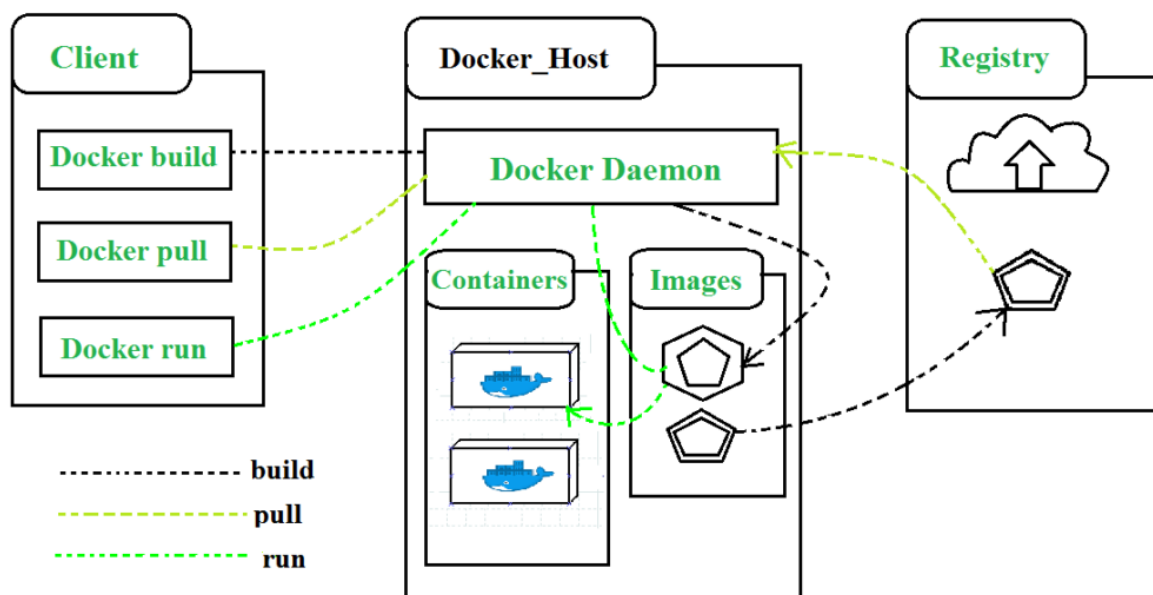
- **Overview of Containerization Technology**
 - **Definition:** Docker is an open-source platform that enables developers to automate the deployment of applications within lightweight, portable containers.
 - **Containerization:** A method of virtualizing an application and its dependencies, allowing it to run consistently across different environments, from development to production.
- **Key Benefits of Using Docker**
 - **Portability:** Applications in Docker containers can run on any system that has Docker installed, regardless of the underlying operating system or infrastructure.
 - **Scalability:** Docker allows for quick scaling of applications by deploying multiple containers as needed.
 - **Isolation:** Each container runs in its own isolated environment, ensuring that dependencies and configurations do not conflict.
 - **Efficiency:** Containers share the host OS kernel, which leads to less overhead compared to traditional VMs, allowing for faster startup and better resource utilization.
 - **Simplified Development and Deployment:** Docker simplifies the development process by enabling consistent environments, which reduces the "it works on my machine" problem.

1.2 Understanding Containerization

- **Differences Between Containers and Traditional Virtualization**
 - **Virtual Machines (VMs):**
 - Each VM runs a complete operating system (guest OS), including a full kernel, which results in higher resource consumption.
 - VMs require a hypervisor to manage them, leading to additional overhead.
 - **Containers:**
 - Containers share the host OS kernel but run isolated user processes. This makes them lightweight and faster to start.
 - There is no need for a hypervisor; instead, container engines like Docker manage the containers directly on the host OS.
- **How Containers Enhance Application Deployment**
 - **Rapid Deployment:** Containers can be started and stopped in seconds, enabling quick deployment cycles.
 - **Consistent Environments:** By encapsulating applications with their dependencies, containers ensure that applications behave the same way in development, testing, and production environments.
 - **Microservices Architecture:** Containers make it easier to adopt microservices architecture by allowing individual services to be developed, deployed, and scaled independently.

1.3 Key Docker Concepts

- **Docker Images**
 - **Definition:** Docker images are read-only templates used to create containers. An image contains the application code, libraries, dependencies, and runtime needed for the application to run.
 - **Layers:** Docker images are built in layers, which helps in efficient storage and reuse of common layers across different images.
- **Containers**
 - **Definition:** A container is a runnable instance of a Docker image. Containers can be started, stopped, moved, and deleted, and they encapsulate the application and its environment.
 - **Lifecycle:** Understanding the lifecycle of a container (created, running, stopped, deleted) is essential for managing them effectively.
- **Dockerfiles**
 - **Definition:** A Dockerfile is a text file that contains instructions for building a Docker image. It defines how the image is created, including the base image, application code, and any required dependencies.
 - **Common Instructions:**
 - FROM: Specifies the base image.
 - RUN: Executes commands in the image.
 - COPY: Copies files into the image.
 - CMD: Specifies the command to run when a container starts.
- **Docker Hub and Container Registries**
 - **Docker Hub:** A public cloud-based repository for sharing and distributing Docker images. Users can pull images from Docker Hub and also push their own images to share with others.
 - **Private Registries:** Organizations may use private registries to host their own images securely. Examples include Google Container Registry and Azure Container Registry.



Chapter 2: Installing and Setting Up Docker

2.1 Installing Docker

- **Step-by-Step Installation on Various Operating Systems**

Windows:

- **System Requirements:** Ensure Windows 10 64-bit (Pro, Enterprise, or Education) or Windows Server 2016/2019.
- **Installation Steps:**
 1. Download Docker Desktop for Windows from the Docker website.
 2. Run the installer and follow the prompts.
 3. Enable WSL 2 (Windows Subsystem for Linux) and install a Linux distribution from the Microsoft Store.
 4. After installation, launch Docker Desktop and follow the setup wizard.

macOS:

- **System Requirements:** macOS Sierra 10.12 or newer.
- **Installation Steps:**
 1. Download Docker Desktop for Mac from the Docker website.
 2. Open the downloaded .dmg file and drag Docker to the Applications folder.
 3. Launch Docker from Applications and follow the setup prompts.

Linux:

- **System Requirements:** Docker supports various distributions (e.g., Ubuntu, CentOS, Debian).
- **Installation Steps** (Ubuntu example):
 1. Update the package index: `sudo apt-get update`.
 2. Install required packages: `sudo apt-get install apt-transport-https ca-certificates curl software-properties-common`.
 3. Add Docker's official GPG key: `curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -`.
 4. Add Docker's repository: `sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"`.
 5. Update the package index again: `sudo apt-get update`.
 6. Install Docker: `sudo apt-get install docker-ce`.
- **Post-Installation Setup and Verification:**
 - **Manage Docker as a non-root user:**
 1. Create the Docker group: `sudo groupadd docker`.
 2. Add your user to the Docker group: `sudo usermod -aG docker $USER`.
 3. Log out and log back in to apply the changes.
 - **Verify Docker Installation:**
 - Run `docker --version` to check the installed version.

- Run `docker run hello-world` to verify that Docker is installed correctly. This command downloads a test image and runs it in a container, displaying a confirmation message.

2.2 Basic Docker Commands

- **Overview of Essential Docker CLI Commands:**
 - **docker run:** Create and start a container from an image.
 - Example: `docker run -d -p 80:80 nginx` (Runs an NGINX web server in detached mode).
 - **docker ps:** List running containers.
 - Example: `docker ps` (Shows active containers).
 - **docker ps -a:** List all containers, including stopped ones.
 - **docker pull:** Download an image from a Docker registry (e.g., Docker Hub).
 - Example: `docker pull ubuntu` (Downloads the latest Ubuntu image).
 - **docker images:** List available images on the local system.
 - **docker rm:** Remove one or more stopped containers.
 - Example: `docker rm <container_id>` (Removes a specified container).
 - **docker rmi:** Remove one or more images.
 - Example: `docker rmi <image_id>` (Removes a specified image).
 - **docker exec:** Execute a command in a running container.
 - Example: `docker exec -it <container_id> bash` (Accesses the container's shell).
- **Hands-on Exercise: Running Your First Container:**
 1. Open a terminal or command prompt.
 2. Run the command: `docker run -d -p 8080:80 nginx`.
 - This command starts an NGINX server and maps port 80 in the container to port 8080 on your host.
 3. Open a web browser and navigate to `http://localhost:8080`. You should see the default NGINX welcome page.

2.3 Configuring Docker Environment

- **Managing Docker Settings:**
 - **Storage:**
 - Understand Docker's storage drivers and how they manage images and containers.
 - Configuring data volumes for persistent storage.
 - Using `docker volume` commands to create and manage volumes.
 - **Network:**
 - Overview of Docker networking modes (bridge, host, overlay).
 - Creating custom networks using `docker network create`.
 - Connecting containers to networks for service communication.
- **Best Practices for Docker Configuration:**
 - **Use Named Volumes:** For persistent data, prefer named volumes over bind mounts to simplify backup and migration.
 - **Limit Resource Usage:** Specify resource limits (CPU and memory) for containers to prevent resource hogging.
 - Example: `docker run --memory="512m" --cpus="1.0" <image>`.

- **Use Environment Variables:** Pass configuration through environment variables instead of hardcoding them in images.
- **Keep Images Small:** Optimize Dockerfiles to reduce image size and improve build time.
- **Regular Updates:** Keep Docker and your images updated to benefit from the latest features and security patches.

Chapter 3: Building Docker Images

3.1 Creating Dockerfiles

- **Writing Effective Dockerfiles:**
 - **Definition:** A Dockerfile is a script containing a series of commands and instructions to create a Docker image.
 - **Best Practices:**
 - Start with a clear base image.
 - Minimize the number of instructions to improve build performance.
 - Use comments to explain non-obvious commands for better maintainability.
- **Instructions and Syntax:**
 - **FROM:** Specifies the base image.
 - Example: FROM ubuntu:20.04 (Sets Ubuntu 20.04 as the base image).
 - **RUN:** Executes commands during the image build process.
 - Example: RUN apt-get update && apt-get install -y curl (Installs curl).
 - **COPY:** Copies files or directories from the host into the image.
 - Example: COPY ./app /usr/src/app (Copies local app directory into the image).
 - **CMD:** Specifies the default command to run when the container starts.
 - Example: CMD ["node", "app.js"] (Runs the Node.js application).
 - **ENTRYPOINT:** Configures a container to run as an executable.
 - Example: ENTRYPOINT ["python"] (Sets Python as the entry point).
 - **EXPOSE:** Informs Docker that the container listens on the specified network ports at runtime.
 - Example: EXPOSE 8080 (Indicates that the application will use port 8080).

3.2 Building Images

- **Using docker build to Create Images from Dockerfiles:**
 - Command syntax: docker build -t <image-name>:<tag> <path-to-Dockerfile>.
 - Example: docker build -t myapp:1.0 . (Builds an image named myapp with the tag 1.0 from the current directory).
 - **Context:** The directory specified at the end of the build command, which includes the Dockerfile and any files needed for the build.
- **Understanding Image Layers and Caching:**
 - **Layers:** Each instruction in a Dockerfile creates a layer in the image, which helps in optimizing storage and reusing unchanged layers during builds.

- **Caching:** Docker caches the layers, so if a layer hasn't changed, Docker uses the cached version instead of rebuilding it. This significantly speeds up build times.
- **Best Practices:**
 - Order instructions to leverage caching effectively. Place frequently changed commands at the end of the Dockerfile.
 - Combine commands where possible to reduce the number of layers.

3.3 Best Practices for Image Optimization

- **Minimizing Image Size and Improving Build Times:**
 - **Use Minimal Base Images:** Choose lightweight base images (e.g., alpine) to reduce size.
 - **Remove Unnecessary Files:** Clean up after installations (e.g., remove package manager caches) using `RUN rm -rf /var/lib/apt/lists/*`.
- **Using Multi-Stage Builds:**
 - **Definition:** Multi-stage builds allow you to use multiple `FROM` statements in a Dockerfile, helping create smaller production images by separating the build environment from the final image.
 - **Example:**

```
Dockerfile
Copy code
# Builder Stage
FROM node:14 AS builder
WORKDIR /app
COPY package.json .
RUN npm install
COPY . .
RUN npm run build

# Production Stage
FROM nginx:alpine
COPY --from=builder /app/build /usr/share/nginx/html
```

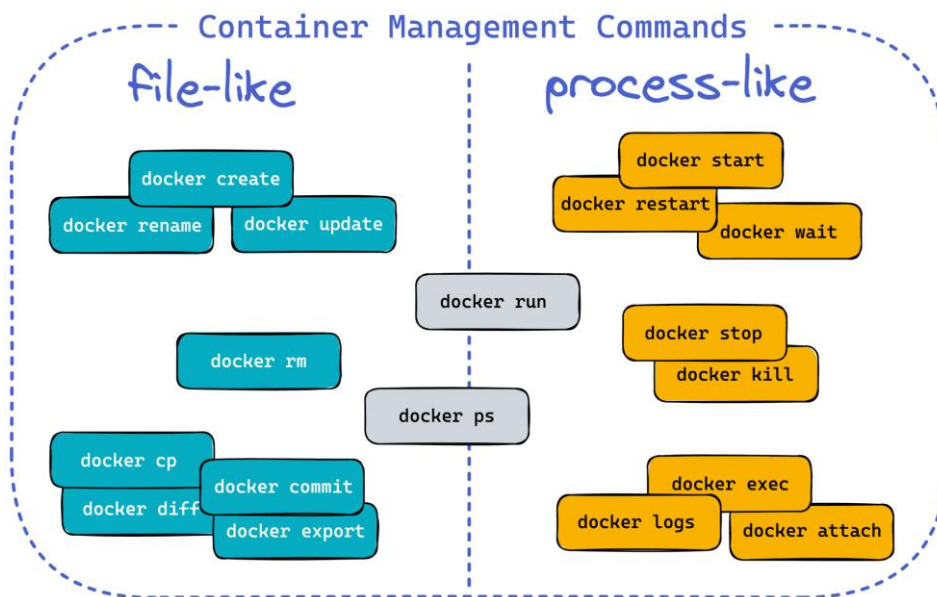
- **Reducing the Number of Layers:**
 - **Combine Commands:** Use `&&` to chain commands in a single `RUN` instruction to minimize the number of layers created.
 - **Use .dockerignore:** Similar to `.gitignore`, this file specifies which files and directories should be excluded from the build context, thus reducing the image size.

Chapter 4: Managing Docker Containers

4.1 Running Containers

- **Starting and Stopping Containers:**
 - **Starting a Container:**
 - Command: `docker run -d --name <container-name> <image-name>`.

- Example: `docker run -d --name my-nginx nginx` (Starts an NGINX container in detached mode).
- **Stopping a Container:**
 - Command: `docker stop <container-name>`.
 - Example: `docker stop my-nginx` (Stops the running NGINX container).
- **Removing a Container:**
 - Command: `docker rm <container-name>`.
 - Example: `docker rm my-nginx` (Removes the stopped container).



- **Understanding Container Lifecycle:**
 - **States:** Containers can be in various states: created, running, paused, stopped, and exited.
 - **Lifecycle Commands:** Familiarize with `docker ps`, `docker ps -a`, and `docker logs <container-name>` to manage and troubleshoot containers.

4.2 Networking in Docker

- **Overview of Docker Networking Concepts:**
 - **Bridge Network:** The default network for containers that are not explicitly connected to any other network.
 - **Host Network:** Containers share the host's network stack, allowing for faster communication but losing isolation.
 - **Overlay Network:** Used for multi-host networking in Docker Swarm, enabling communication between containers on different hosts.
- **Configuring Container Networking for Service Communication:**
 - **Creating Custom Networks:**
 - Command: `docker network create <network-name>`.
 - Example: `docker network create my-network` (Creates a custom bridge network).
 - **Connecting Containers to Networks:**

- Command: `docker network connect <network-name> <container-name>`.
- Command: `docker run --network <network-name> <image-name>` to start a container on a specific network.

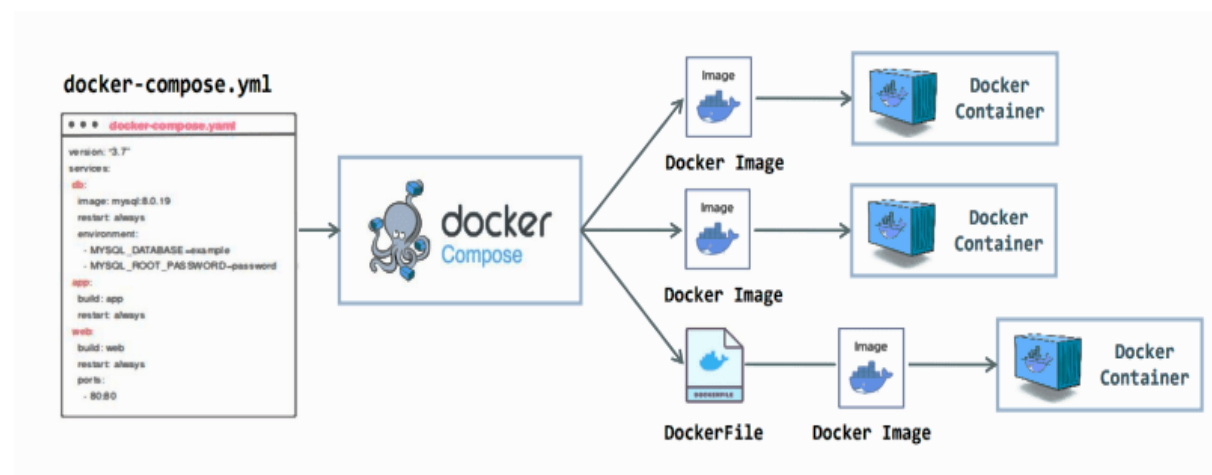
4.3 Volume Management

- **Using Volumes for Persistent Data Storage:**
 - **Definition:** Volumes are a way to persist data generated by and used by Docker containers, ensuring data isn't lost when containers stop or are removed.
 - **Creating a Volume:**
 - Command: `docker volume create <volume-name>`.
 - Example: `docker volume create my-data` (Creates a named volume).
 - **Mounting Volumes:**
 - Command: `docker run -v <volume-name>:/path/in/container <image-name>`.
 - Example: `docker run -d -v my-data:/data nginx` (Mounts the my-data volume at /data in the NGINX container).
- **Best Practices for Managing Data in Docker Containers:**
 - **Use Named Volumes:** Prefer named volumes over bind mounts for better portability and management.
 - **Regular Backups:** Implement a backup strategy for volume data to avoid data loss.
 - **Manage Volume Lifecycles:** Use `docker volume ls`, `docker volume inspect`, and `docker volume rm` to manage volumes efficiently.

Chapter 5: Docker Compose for Multi-Container Applications

5.1 Introduction to Docker Compose

- **Overview of Docker Compose and Its Use Cases:**
 - **Definition:** Docker Compose is a tool for defining and running multi-container Docker applications using a single configuration file.



- **Use Cases:**
 - Developing applications with multiple services (e.g., web app, database, cache).
 - Simplifying configuration management for multi-container environments.
 - Facilitating local development and testing.
- **Key Concepts:**
 - **Services:** Individual containers that perform specific functions (e.g., web server, database).
 - **Networks:** Allows services to communicate with each other. Docker Compose creates a default network for the services defined in the docker-compose.yml file.
 - **Volumes:** Persist data generated by and used by the services, ensuring data integrity and availability.

5.2 Defining Multi-Container Applications

- **Writing docker-compose.yml Files:**
 - **Basic Structure:**

```
yaml
Copy code
version: '3'
services:
  web:
    image: nginx
    ports:
      - "8080:80"
  db:
    image: mysql
    environment:
      MYSQL_ROOT_PASSWORD: example
```

- **Version:** Defines the version of the Compose file format.
 - **Services:** Lists the services to be deployed.
 - **Ports:** Maps container ports to host ports.
 - **Environment Variables:** Passes environment variables to services (e.g., setting database passwords).
- **Hands-On Exercise: Deploying a Multi-Container Application:**
 - **Example Application:** Create a simple web application using NGINX and a MySQL database.
 - **Steps:**
 1. Create a docker-compose.yml file with the required services.
 2. Use docker-compose up to start the application.
 3. Access the web application via a browser and verify database connectivity.

5.3 Managing and Scaling Applications

- **Using Docker Compose Commands to Manage Applications:**

- **Basic Commands:**
 - docker-compose up: Starts the defined services.
 - docker-compose down: Stops and removes the services.
 - docker-compose ps: Lists the running services.
 - docker-compose logs: Displays logs for the services.
- **Scaling Services with Docker Compose:**
 - **Scaling:** Allows you to run multiple instances of a service for load balancing and high availability.
 - **Command:**
 - Use docker-compose up --scale <service-name>=<number> to scale services.
 - Example: docker-compose up --scale web=3 (Runs three instances of the web service).
 - **Best Practices:** Consider the implications of scaling on service communication and resource utilization.

Chapter 6: Deploying Docker Containers

6.1 Deployment Strategies

- **Overview of Deployment Options:**
 - **Local Deployment:** Running containers on a local machine for development and testing.
 - **Cloud Deployment:** Utilizing cloud platforms (e.g., AWS, Azure, Google Cloud) to deploy Docker containers for production.
 - **Orchestration:** Using orchestration tools (e.g., Docker Swarm, Kubernetes) to manage the deployment, scaling, and operation of containerized applications.
- **Best Practices for Deploying Containers:**
 - **Environment Consistency:** Ensure consistency between development, staging, and production environments to minimize issues.
 - **Configuration Management:** Use environment variables or external configuration files for sensitive data and settings.
 - **Monitoring and Logging:** Implement monitoring and logging solutions for proactive issue detection and troubleshooting.

6.2 Using Docker Swarm

- **Introduction to Docker Swarm for Container Orchestration:**
 - **Definition:** Docker Swarm is a native clustering and orchestration tool for Docker, enabling the management of a cluster of Docker nodes as a single virtual system.
 - **Features:** Load balancing, service discovery, scaling, and rolling updates.
- **Configuring a Swarm Cluster and Deploying Services:**
 - **Creating a Swarm:**
 - Command: docker swarm init to initialize a Swarm on the current node.
 - Command: docker swarm join to add other nodes to the Swarm.
 - **Deploying Services:**

- Command: `docker service create --name <service-name> <image-name>`.
 - Example: `docker service create --name web nginx` (Deploys an NGINX service in the Swarm).
- **Managing Services:**
 - Use commands like `docker service ls`, `docker service scale`, and `docker service rm` to manage deployed services.

6.3 Using Kubernetes with Docker

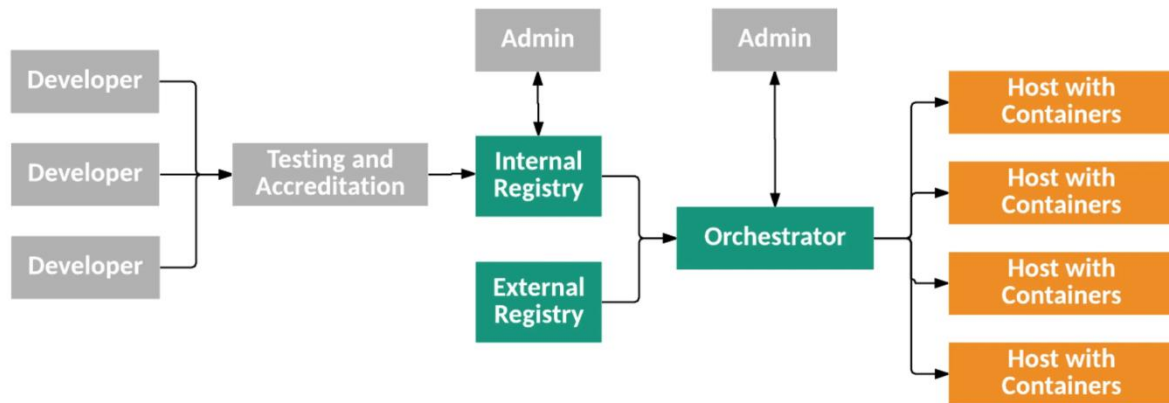
- **Overview of Kubernetes as a Container Orchestration Platform:**
 - **Definition:** Kubernetes is an open-source platform for automating the deployment, scaling, and management of containerized applications.
 - **Key Concepts:**
 - **Pods:** The smallest deployable units in Kubernetes, which can contain one or more containers.
 - **Deployments:** Manage the desired state for Pods, allowing for rolling updates and scaling.
 - **Services:** Define a logical set of Pods and enable communication between them.
- **Deploying Docker Containers in Kubernetes (Basic Concepts and Setup):**
 - **Setting Up Kubernetes:**
 - Use Minikube for local Kubernetes development or managed services (e.g., GKE, EKS) for cloud deployments.
 - **Creating a Deployment:**
 - Example deployment.yaml:

```

yaml
Copy code
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web
spec:
  replicas: 3
  selector:
    matchLabels:
      app: web
  template:
    metadata:
      labels:
        app: web
    spec:
      containers:
      - name: nginx
        image: nginx
      ports:
      - containerPort: 80

```

- Command: `kubectl apply -f deployment.yaml` to deploy the application.
- **Exposing the Application:**
 - Command: `kubectl expose deployment web --type=LoadBalancer --port=80` to create a service for external access.



Chapter 7: Monitoring and Logging

7.1 Monitoring Docker Containers

- **Tools and Techniques for Monitoring Container Performance:**
 - **Prometheus:**
 - **Definition:** An open-source monitoring and alerting toolkit designed for reliability and scalability.
 - **Features:** Pull-based data collection, flexible querying language (PromQL), and powerful alerting capabilities.
 - **Grafana:**
 - **Definition:** An open-source visualization and analytics platform.
 - **Integration:** Works seamlessly with Prometheus for visualizing metrics and creating dashboards.
 - **Hands-on Exercise:** Set up Prometheus and Grafana to monitor Docker containers.
- **Understanding Resource Usage and Metrics:**
 - **Key Metrics:**
 - CPU usage: Monitor how much CPU a container consumes.
 - Memory usage: Keep track of memory allocation and usage.
 - Network I/O: Measure incoming and outgoing traffic for containers.
 - Disk I/O: Monitor read/write operations on volumes.
 - **Visualizations:** Create dashboards in Grafana to represent container performance metrics.


7.2 Logging in Docker

- **Configuring Logging Drivers in Docker:**
 - **Overview of Docker Logging Drivers:**
 - **Built-in Drivers:** json-file, syslog, journald, fluentd, gelf, etc.

- **Configuration:** Example of configuring logging driver in docker run:

```
bash
Copy code
docker run --log-driver=syslog <image>
```

- **Choosing the Right Driver:** Selecting a logging driver based on the application and infrastructure needs.
- **Best Practices for Managing Container Logs:**
 - **Log Rotation:** Implement log rotation to prevent disk space issues.
 - **Structured Logging:** Use structured log formats (JSON) for easier parsing and analysis.
 - **Retention Policies:** Define retention policies to manage log data lifecycle.



7.3 Centralized Logging Solutions

- **Integrating Docker with Centralized Logging Tools:**
 - **ELK Stack:**
 - **Components:** Elasticsearch, Logstash, and Kibana.
 - **Functionality:** Collects, stores, and visualizes log data from multiple sources.
 - **Hands-on Exercise:** Set up an ELK stack to collect and visualize Docker logs.
 - **Fluentd:**
 - **Definition:** An open-source data collector for unified logging.
 - **Integration:** Use Fluentd to collect logs from Docker containers and send them to various backends (e.g., Elasticsearch).
- **Analyzing Logs for Troubleshooting and Performance Tuning:**
 - **Log Analysis Techniques:**
 - Search and filter logs to identify issues.
 - Correlate logs from different services for comprehensive insights.
 - **Performance Tuning:** Use log data to identify bottlenecks and optimize application performance.

Chapter 8: Security Best Practices for Docker

8.1 Understanding Docker Security Risks

- **Common Vulnerabilities and Threats in Containerized Environments:**
 - **Vulnerabilities:** Outdated images, misconfigurations, and insecure code.
 - **Threats:** Unauthorized access, data breaches, and denial of service attacks.
- **Importance of Security in the Container Lifecycle:**
 - **Security throughout the Lifecycle:** Emphasizing the need for security at every stage—from development to deployment and maintenance.

8.2 Implementing Security Measures

- **Best Practices for Securing Docker Containers:**
 - **User Permissions:** Run containers as non-root users to minimize security risks.
 - **Image Scanning:** Regularly scan images for vulnerabilities using tools like Clair or Trivy.
 - **Minimize Attack Surface:** Use minimal base images and remove unnecessary packages.
- **Network Security Considerations for Containerized Applications:**
 - **Network Segmentation:** Isolate containers within networks to limit exposure.
 - **Firewall Rules:** Configure firewalls to restrict incoming and outgoing traffic.
 - **Encryption:** Use TLS for securing communication between containers.

8.3 Compliance and Governance

- **Tools and Frameworks for Ensuring Compliance in Docker Environments:**
 - **Compliance Frameworks:** Familiarize with frameworks like NIST, ISO 27001, and CIS benchmarks for container security.
 - **Compliance Tools:** Utilize tools such as Aqua Security, Sysdig, and Twistlock for compliance checks and reporting.
- **Continuous Security Practices in DevOps:**
 - **Integration with CI/CD:** Incorporate security checks into the CI/CD pipeline to catch vulnerabilities early.
 - **Automated Monitoring:** Implement continuous monitoring solutions to detect security issues in real-time.

Summary and Key Takeaways

- **Recap of Key Concepts:**
 - Summarize the major topics covered in the course, highlighting the importance of monitoring, logging, and security in Docker.
- **Importance of Docker in Modern Development:**
 - Discuss how Docker facilitates DevOps practices, improves application deployment, and enhances software delivery processes.
- **Future Trends in Containerization:**
 - Explore emerging trends such as serverless containers, container-native security, and advancements in orchestration technologies.